

DSBA 6188 - Homework 2

Manjinder Sandhu

Task: The task for this assignment is to develop a text classifier that can filter through posts and comments on Reddit Cooking Threads and categorize content relevant to recipes. It will not include unrelated discussions and noise. By implementing this classifier, stakeholders will be able to specifically focus on recipe topics and gather insight into cooking trends on Reddit.

Approach and Reasoning

When we started this assignment we had multiple people label a dataset called “cooking_eval_reviewed2.jsonl” to create a golden evaluation dataset. By having multiple people and myself label this dataset it ensures consistency and prevents any discrepancies that can affect the dataset. This approach reduced errors in the dataset and created a more accurate and trustworthy evaluation. We are going to use this “golden dataset” to evaluate the training model. We used “ngrok” to accomplish this.

In the appendix, there will be the inter-annotator agreement that my teammates came up with. The inter-annotator agreement shows the level of agreement between the annotators by seeing how often the same label was annotated on the same data. This is very important because it allows us to understand the quality of the annotations and ensure the labeled dataset is reliable. ChatGPT was also used to help with coding, understanding concepts, and fixing errors.

The training dataset that I used to label my data was called “homework2_training”. This dataset had 5,000 unlabeled posts and comments from Reddit. I labeled 500 posts/comments then I did 2 experiments. The first experiment was with Prodigy Train where on average my evaluation score was 0.74. The second experiment I did was with spacy train and used a based model called “en_core_web_lg”, my evaluation score was on average around 0.925. This was done only on 10 percent of the data. I wanted a way where I could label all 5,000 posts/comments without spending too much time. I used a concept called Weakly labels.

Weakly labels is a technique in spaCy where instead of manually labeling your annotations you can label them using heuristic rules, distance supervision, and bootstrapping. These methods allow you to assign labels to large datasets with minimal human effort. Weakly labeling may not be as precise or reliable as manually labeling a dataset. There are several techniques to implement a Weakly label:

1. **Bootstrapping:**
 - a. This is where you manually label a small amount of the data and then use a model train on the data to label the rest. I used this method.
2. **Heuristic Rules:**

- a. This is where you establish patterns and rules to automatically label your data. For example, in this dataset, any keywords related to cooking might be weakly labeled as a recipe.
3. **Distance Supervision:**
 - a. This is where you use metadata or other information from the database to assign labels. You are using external sources to label the data.
4. **Crowdsourcing:**
 - a. This is where you get a team and everyone labels a dataset.

When I applied Weakly label I broke down each step instead of writing one code that does everything. I did this to better apply the concept and fully understand it.

Step 1:

I first had to get my 500 labeled data and export it out of the database. I did this by

```
prodigy db-out homework2_manual > exported_data.jsonl
```

After getting those 500 labels, I had to format them in a way that spaCy could read them. I want spaCy to read my annotations so I can train my model. I did this by using the “readSp.py” code and the output file was called “processed_data.jsonl”. The processed_data.jsonl files had the labels “RELEVANT” or “NOT_RELEVANT” for each jsonl line.

Step 2:

After formatting the jsonl file for spaCy. I used a Python script called “score_records.py” that trained a text classification model to tell the difference between “RELEVANT” or “NOT_RELEVANT”. It initialized a blank English spaCy model and added a text classification pipeline to it. Then it used stochastic gradient descent (SGD) to train it. Once it is done training it is saved to the local disk for future text classification. You can also find this in the file called “my_trained_model”.

Step 3:

In “codethatLabel.py”, we load the pre-trained classification model using spaCy and we iterate through each JSONL line applying the model to predict the labels. We add the original unlabeled files as the input file. It will give a “RELEVANT” and “NOT_RELEVANT” score. The output for this is “classified_data.jsonl”.

Step 4:

The last step is to transform the classified data to another format based on the threshold. We can do this by running “finalstep.py”. The threshold is 0.5. Based on the threshold it will assign “accept” or “reject” to the record. The code generates an output record for each record. The output file for this is called “homework2_trainComplete.jsonl” with 5000 records annotated.

I will assess my new training data again against my golden dataset. The third experiment was with Prodigy Train where on average my evaluation score was 0.75. The fourth experiment I did was with the spaCy train and used a based model called “en_core_web_lg”, my evaluation score was on average around 0.81. There was a slight difference between training 500 vs 5,000 on the evaluation. I believe that the 500 labeled data were the most straightforward to distinguish. When I used 100 percent of the data, I received a more diverse range of examples that had more challenging data. To fix this I can perform data augmentation and regularization techniques.

There are several files that I did not include in the final GitHub Repositories like the venv, output/experiment-2 because those files were over 500 MB. They are in the git ignore.

Annotated Guidelines for Classification on Recipes

Task: The goal of these guidelines is to provide instructions for labeling “RELEVANT” or “NOT_RELEVANT” for specific comments/posts on Reddit Cooking Threads.

Definition:

RELEVANT: This is a post/comment on Reddit Cooking Threads that is related to recipes. It can include instructions for preparing a dish or discussing ingredients for a recipe.

NOT_RELEVANT: This is a post/comment on Reddit Cooking Threads that isn’t related to recipes. It doesn’t contain any details on how to prepare a dish or doesn’t have any food-related content.

Examples:

✓ It includes 5 pounds of flour and 6 tablespoons of olive oil.

✗ It has 4 knives and 2 blenders

✓ I like to use flatbread. Naan or Pati. Ensure it is whole grain. For vegetables use tomato, asparagus, artichokes, and peas.

✗ I like to cook my flatbread on a stove because it says it in a recipe book.

There are no recipes in these comments, but they are related to discussing ingredients in a dish so the green is RELEVANT while the red is NOT_RELEVANT. The red doesn’t help add anything meaningful to the dish discussion or come up with new recipes.

Proposal For Recipe Classifier For Different Cuisine Types

Objective: The goal for this is to create a machine-learning model that can classify recipes based on cuisine type. Some of the cuisines that we may include are Mexican, Italian, Indian, Mediterranean, and many more.

Approach:

Data Collection: We want to gather different datasets that have different recipes from various cuisine types. We are going to look for publicly available datasets or APIs to get this information.

Data Preprocessing: In this step, we will tokenize the recipe into words for further processing. It is important to conduct data augmentation techniques to diversify the dataset. We can perform standardization to remove irrelevant information.

Model Training: In this step, we will use spaCy to implement a multi-class text classification model. It is also good to fine-tune the pre-trained model and I recommend that we use “en_core_web_lg”.

Annotations: In this stage, we want to annotate the data efficiently so we will use Prodigy. We can get a group of people to annotate different cuisine types they see in the dataset.

Evaluation: Then we have to assess the model performance by looking at the F1-score. We can also use cross-validation to ensure that the model can perform well on unseen data. If the model performs low then we can go back and review the annotations to further improve the model.

Timeline: It will take our company around 8 months to complete this project. The 8 months will factor in delays and other time-consuming factors that may affect the project.

Teams:

- NLP Engineers
- Data Scientist
- Project Manager

Conclusion:

We are going to provide a high-quality machine-learning model that classifies recipes based on cuisine type. We will use tools such as Prodigy and spaCy. This project will take around 8 months to complete.

Appendix:

Inter-Annotator Agreement:

```
i Using 4 annotator IDs: cooking_eval-Tyler, cooking_eval-alex,  
cooking_eval-Manjinder, cooking_eval-"dagim"
```

```
i Annotation Statistics
```

Attribute	Value
-----	-----
Examples	800
Categories	2
Co-Incident Examples*	200
Single Annotation Examples	0
Annotators	4
Avg. Annotations per Example	4.0

```
* (>1 annotation)
```

```
i Agreement Statistics
```

Statistic	Value
-----	-----
Percent (Simple) Agreement	0.7067
Krippendorff's Alpha	0.4075
Gwet's AC2	0.4197

Experiment 1: 500 Prodigy Train

===== Training pipeline =====

Components: textcat_multilabel

Merging training and evaluation data for 1 components

- [textcat_multilabel] Training: 500 | Evaluation: 200 (from datasets)

Training: 500 | Evaluation: 200

Labels: textcat_multilabel (2)

i Pipeline: ['textcat_multilabel']

i Initial learn rate: 0.001

E	#	LOSS TEXTC...	CATS_SCORE	SCORE
0	0	0.25	34.88	0.35
0	200	46.12	69.38	0.69
2	400	27.12	75.93	0.76
3	600	18.18	73.69	0.74
4	800	13.67	74.39	0.74
6	1000	9.80	73.92	0.74
8	1200	7.61	74.09	0.74
11	1400	5.59	73.72	0.74
14	1600	4.41	73.89	0.74
18	1800	3.39	73.87	0.74
23	2000	2.57	73.94	0.74

Experiment 2: 500 spaCy

To use this data for training with spaCy, you can run:

```
python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
```

```
(venv) (venv) Manjinders-MacBook-Pro:homework2 manjinders$ python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
```

i No output directory provided

i Using CPU

===== Initializing pipeline =====

✓ Initialized pipeline

===== Training pipeline =====

i Pipeline: ['tok2vec', 'tagger', 'parser', 'attribute_ruler',

'lemmatizer', 'ner', 'textcat_multilabel']

i Frozen components: ['tagger', 'parser', 'attribute_ruler',

'lemmatizer', 'ner']

i Initial learn rate: 0.001

E	#	LOSS TOK2VEC	LOSS TEXTC...	CATS_SCORE	SPEED	SCORE
0	0	0.02	0.22	43.41	6023.70	0.43
6	1000	38.20	86.59	91.11	7377.56	0.91
23	2000	21.64	4.06	88.72	7038.52	0.89
60	3000	28.00	1.04	88.01	7410.59	0.88
100	4000	47.29	1.20	89.06	7822.06	0.89
140	5000	0.02	0.00	91.66	7321.56	0.92
180	6000	0.00	0.00	91.25	6758.01	0.91
221	7000	0.00	0.00	91.18	7797.50	0.91
261	8000	0.00	0.00	91.32	7691.92	0.91
301	9000	0.00	0.00	91.78	7751.39	0.92
341	10000	0.00	0.00	91.60	7629.94	0.92
381	11000	0.00	0.00	91.69	7598.87	0.92
421	12000	0.00	0.00	91.91	7440.61	0.92
462	13000	0.00	0.00	91.82	7333.26	0.92
502	14000	0.00	0.00	91.86	6927.00	0.92
542	15000	0.00	0.00	92.05	7701.79	0.92
582	16000	0.00	0.00	92.19	7587.03	0.92
622	17000	0.00	0.00	92.37	7479.82	0.92
662	18000	0.00	0.00	92.52	7427.78	0.93
703	19000	0.00	0.00	92.64	7509.34	0.93
743	20000	0.00	0.00	92.75	7500.18	0.93
783	21000	0.00	0.00	92.88	7487.14	0.93
824	22000	0.00	0.00	93.02	7453.42	0.93
864	23000	0.00	0.00	93.13	7510.56	0.93
904	24000	0.00	0.00	93.18	7362.93	0.93
944	25000	0.00	0.00	93.15	7459.23	0.93
985	26000	0.00	0.00	93.16	7430.85	0.93
1025	27000	0.00	0.00	93.13	7595.56	0.93
1065	28000	0.00	0.00	93.08	7432.19	0.93
1105	29000	0.00	0.00	93.11	7498.57	0.93

Experiment 3: 5000 Prodigy Train

```

[textcat_multilabel] Training: 5000 | Evaluation: 200 (from datasets)
Training: 5000 | Evaluation: 200
Labels: textcat_multilabel (2)
i Pipeline: ['textcat_multilabel']
i Initial learn rate: 0.001
E      #      LOSS TEXTC...  CATS_SCORE  SCORE
-----
0      0      0.25      32.27      0.32
0      200     48.93      54.73      0.55
0      400     34.34      65.12      0.65
0      600     35.86      70.19      0.70
0      800     33.11      72.95      0.73
0     1000     29.71      73.90      0.74
0     1200     27.21      73.30      0.73
1     1400     23.50      72.37      0.72
1     1600     21.05      72.91      0.73
1     1800     20.14      73.42      0.73
2     2000     18.57      73.68      0.74
3     2200     16.78      74.70      0.75
3     2400     14.63      74.31      0.74
4     2600     13.19      74.85      0.75
5     2800     12.20      75.02      0.75
6     3000     11.32      75.53      0.76
7     3200     10.05      74.95      0.75
7     3400      8.89      75.09      0.75
8     3600      8.72      75.09      0.75
9     3800      7.60      74.91      0.75
10    4000      7.26      75.41      0.75
11    4200      6.62      75.38      0.75
12    4400      6.12      75.04      0.75
12    4600      6.05      75.31      0.75

```

Experiment 4: 5000 spaCy

```

===== Generating config =====
i Auto-generating config with spaCy
i Using config from base model
✓ Generated training config

===== Generating cached label data =====
✓ Saving label data for component 'tagger'
corpus/labels/tagger.json
✓ Saving label data for component 'parser'
corpus/labels/parser.json
✓ Saving label data for component 'ner'
corpus/labels/ner.json
✓ Saving label data for component 'textcat_multilabel'
corpus/labels/textcat_multilabel.json

===== Finalizing export =====
✓ Saved training config
corpus/config.cfg

To use this data for training with spaCy, you can run:
python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
(venv) (venv) Manjinders-MacBook-Pro:homework2 manjinder$ python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
i No output directory provided
i Using CPU

===== Initializing pipeline =====
✓ Initialized pipeline

===== Training pipeline =====
i Pipeline: ['tok2vec', 'tagger', 'parser', 'attribute_ruler',
'lemmatizer', 'ner', 'textcat_multilabel']
i Frozen components: ['tagger', 'parser', 'attribute_ruler',
'lemmatizer', 'ner']
i Initial learn rate: 0.001
E      #      LOSS TOK2VEC  LOSS TEXTC...  CATS_SCORE  SPEED  SCORE
-----
0      0      0.09      0.29      40.33  4139.42  0.48
0     1000     34.09     179.79     87.48  7990.48  0.87
2     2000     51.11     115.09     84.23  7915.07  0.84
6     3000     99.28      37.74     81.08  6832.15  0.81
10    4000     85.82     13.12     80.13  6611.60  0.80
14    5000     86.32      7.52     80.76  7223.10  0.81
18    6000     78.08      4.43     80.30  7075.75  0.80

```

Commands to run this file:

WARNING: Ensure that you have a Prodigy License. All these Command Lines were done in the terminal.

```
source venv/bin/activate
```

```
python3 -m prodigy textcat.manual homework2_manual  
data/homework2_train.jsonl --label RELEVANT,NOT_RELEVANT
```

```
(venv) Manjinders-MacBook-Pro:homework2 manjinder$ source venv/bin/activate  
○ (venv) (venv) Manjinders-MacBook-Pro:homework2 manjinder$ python3 -m prodigy textcat.manual homework2_manual data/homework2_train.jsonl --label RELEVANT,NOT_RELEVANT  
Using 2 label(s): RELEVANT, NOT_RELEVANT  
★ Starting the web server at http://localhost:8080 ...  
Open the app in your browser and start annotating!
```

```
python3 -m prodigy train --textcat-multilabel  
homework2_manual,eval:hmwk2-eval ./output/experiment-1
```

```
===== Training pipeline =====  
Components: textcat_multilabel  
Merging training and evaluation data for 1 components  
- [textcat_multilabel] Training: 500 | Evaluation: 200 (from datasets)  
Training: 500 | Evaluation: 200  
Labels: textcat_multilabel (2)  
i Pipeline: ['textcat_multilabel']  
i Initial learn rate: 0.001  
E   #   LOSS TEXTC...   CATS_SCORE   SCORE  
---  ---  
0     0         0.25         34.88        0.35  
0    200        46.12        69.38        0.69  
2    400        27.12        75.93        0.76  
3    600        18.18        73.69        0.74  
4    800        13.67        74.39        0.74  
6   1000         9.80        73.92        0.74  
8   1200         7.61        74.09        0.74  
11  1400         5.59        73.72        0.74  
14  1600         4.41        73.89        0.74  
18  1800         3.39        73.87        0.74  
23  2000         2.57        73.94        0.74
```

```
python3 -m prodigy data-to-spacy --textcat-multilabel  
homework2_manual,eval:hmwk2-eval ./corpus --base-model  
en_core_web_lg
```

```
python -m spacy train corpus/config.cfg --paths.train  
corpus/train.spacy --paths.dev corpus/dev.spacy
```


python3 readSp.py

```
1 import jsonlines
2
3 # Path to your dataset file
4 dataset_file = "exported_data.jsonl"
5
6 # Path to the output file
7 output_file = "processed_data.jsonl"
8
9 # Open the JSONL file and extract text and labels
10 try:
11     with jsonlines.open(dataset_file) as reader, jsonlines.open(output_file, mode='w') as writer:
12         for obj in reader:
13             text = obj.get("text")
14             label = obj.get("accept", []).get(0) # Get the first accepted label if available
15             if text and label:
16                 writer.write({"text": text, "label": label})
17             else:
18                 print("Warning: Text or label missing in the JSON object.")
19         print("Processing completed. Output written to:", output_file)
20 except Exception as e:
21     print("Error:", e)
22
```

python3 score_records.py

```
score_records.py > ...
1 import spacy
2 from spacy.training import Example
3 import jsonlines
4 import random
5 # Load a blank English model
6 nlp = spacy.blank("en")
7 # Add text classification pipeline to the model
8 textcat = nlp.add_pipe('textcat_multilabel', last=True)
9 textcat.add_label("RELEVANT")
10 textcat.add_label("NOT_RELEVANT")
11
12 # Path to the processed data file
13 processed_data_file = "processed_data.jsonl"
14
15 # Open the JSONL file and extract text and labels
16 with jsonlines.open(processed_data_file) as reader:
17     processed_data = list(reader)
18
19 # Convert processed data to spaCy format
20 spacy_train_data = []
21 for obj in processed_data:
22     text = obj["text"]
23     label = {"RELEVANT": obj["label"] == "RELEVANT", "NOT_RELEVANT": obj["label"] == "NOT_RELEVANT"}
24     spacy_train_data.append(Example.from_dict(nlp.make_doc(text), {"cats": label}))
25
26 # Initialize the model and get the optimizer
27 optimizer = nlp.initialize()
28
29 # Train the text classification model
30 n_iter = 10
31 for i in range(n_iter):
32     spacy.util.fix_random_seed(1)
33     random.shuffle(spacy_train_data)
34     losses = {}
35     for batch in spacy.util.minibatch(spacy_train_data, size=8):
36         nlp.update(batch, losses=losses, sgd=optimizer)
37     print("Iteration:", i, "Losses:", losses)
38
39 # Save the trained model
40 output_dir = "./my_trained_model"
41 nlp.to_disk(output_dir)
```

python3 codethatLabel.py

```
1 import spacy
2 import jsonlines
3
4 # Load the trained model
5 model_path = "./my_trained_model"
6 nlp = spacy.load(model_path)
7
8 # Load the unlabeled data
9 unlabeled_data_file = "data/homework2_train.jsonl"
10
11 # Open the JSONL file and classify each record
12 classified_data = []
13 with jsonlines.open(unlabeled_data_file) as reader:
14     for record in reader:
15         text = record["text"]
16         doc = nlp(text)
17         predicted_labels = doc.cats
18         classified_data.append({"text": text, "predicted_labels": predicted_labels})
19
20 # Optionally, you can save the classified data to a file or process it further
21 output_file = "classified_data.jsonl"
22 with jsonlines.open(output_file, mode="w") as writer:
23     writer.write_all(classified_data)
```

prodigy db-in homework2_manual2 homework_trainComplete.jsonl

```
python3 -m prodigy textcat.manual homework2_manual2
homework2_trainComplete.jsonl --label RELEVANT,NOT_RELEVANT
--exclusive
```

```

[textcat_multilabel] Training: 5000 | Evaluation: 200 (from datasets)
Training: 5000 | Evaluation: 200
Labels: textcat_multilabel (2)
i Pipeline: ['textcat_multilabel']
i Initial learn rate: 0.001

```

E	#	LOSS TEXTC...	CATS_SCORE	SCORE
0	0	0.25	32.27	0.32
0	200	48.93	54.73	0.55
0	400	34.34	65.12	0.65
0	600	35.86	70.19	0.70
0	800	33.11	72.95	0.73
0	1000	29.71	73.90	0.74
0	1200	27.21	73.30	0.73
1	1400	23.50	72.37	0.72
1	1600	21.05	72.91	0.73
1	1800	20.14	73.42	0.73
2	2000	18.57	73.68	0.74
3	2200	16.78	74.70	0.75
3	2400	14.63	74.31	0.74
4	2600	13.19	74.85	0.75
5	2800	12.20	75.02	0.75
6	3000	11.32	75.53	0.76
7	3200	10.05	74.95	0.75
7	3400	8.89	75.09	0.75
8	3600	8.72	75.09	0.75
9	3800	7.60	74.91	0.75
10	4000	7.26	75.41	0.75
11	4200	6.62	75.38	0.75
12	4400	6.12	75.04	0.75
12	4600	6.05	75.31	0.75

```

python3 -m prodigy data-to-spacy --textcat-multilabel
homework2_manual2,eval:hmwk2-eval ./corpus --base-model
en_core_web_lg

```

```

python -m spacy train corpus/config.cfg --paths.train
corpus/train.spacy --paths.dev corpus/dev.spacy

```

```

===== Generating config =====
i Auto-generating config with spaCy
i Using config from base model
✓ Generated training config

===== Generating cached label data =====
✓ Saving label data for component 'tagger'
corpus/labels/tagger.json
✓ Saving label data for component 'parser'
corpus/labels/parser.json
✓ Saving label data for component 'ner'
corpus/labels/ner.json
✓ Saving label data for component 'textcat_multilabel'
corpus/labels/textcat_multilabel.json

===== Finalizing export =====
✓ Saved training config
corpus/config.cfg

To use this data for training with spaCy, you can run:
python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
(venv) (venv) Manjinders-MacBook-Pro:homework2 manjinder$ python -m spacy train corpus/config.cfg --paths.train corpus/train.spacy --paths.dev corpus/dev.spacy
i No output directory provided
i Using CPU

===== Initializing pipeline =====
✓ Initialized pipeline

===== Training pipeline =====
i Pipeline: ['tok2vec', 'tagger', 'parser', 'attribute_ruler',
'lemmatizer', 'ner', 'textcat_multilabel']
i Frozen components: ['tagger', 'parser', 'attribute_ruler',
'lemmatizer', 'ner']
i Initial learn rate: 0.001
E #      LOSS TOK2VEC  LOSS TEXTC...  CATS_SCORE  SPEED  SCORE
--
0 0      0.09      0.29      40.33  4139.42  0.40
0 1000    34.09      179.79      87.48  7990.48  0.87
2 2000    51.11      115.09      84.23  7915.07  0.84
6 3000    99.28      37.74      81.08  6832.15  0.81
10 4000    85.82      13.12      80.13  6611.60  0.80
14 5000    86.32      7.52      80.76  7223.10  0.81
18 6000    78.08      4.43      80.30  7075.75  0.80

```